

# Crystal Ball Gazing: Mathematical models for chasing returns

Thomas P. Steele

June 2, 2025

# Contents

<b>1</b>	<b>Mathematical Tools</b>	<b>2</b>
1.1	Stochastic Models and Inference . . . . .	2
1.1.1	Example: Linear Regression . . . . .	4
1.2	Bayes Theorem and Bayesian Inference . . . . .	5
1.2.1	Bayesian Linear Regression . . . . .	6
<b>2</b>	<b>Financial Random Walks</b>	<b>9</b>
2.1	Discrete Time Continuous Space Financial Random Walks . . . . .	9
2.1.1	Time Series Models . . . . .	10
2.2	Continuous Time Continuous Space Financial Random Walks . . . . .	10

# Chapter 1

## Mathematical Tools

### 1.1 Stochastic Models and Inference

A model in general takes some input and gives some output,

$$f(x) = y$$

A stochastic model is the same, however, every single time we evaluate the model we get a different answer.

To get a feel for stochastic modelling, we will start with a basic stochastic model. Suppose we are studying some phenomena and collect a bunch of observations. We will call the set of observations  $Y$  and each observation  $y_i$ . Our goal is to make a stochastic model which mimics our set of observations,

$$y_i = \mu + \varepsilon_i$$

you can think of  $\varepsilon$  as the *error*, so our model is that each of our observations has a value  $\mu$  plus some error. For this type of model we must make an assumption on how our error is distributed, the easiest and often best assumption is that our error is normally distributed with mean 0 and some standard deviation  $\sigma$ ,  $\varepsilon_i \sim \mathcal{N}(0, \sigma)$ . Our model then becomes,

$$y_i \sim \mathcal{N}(\mu, \sigma),$$

meaning, our observations are randomly distributed by the normal distribution with mean  $\mu$  and standard deviation  $\sigma$ . To actually learn anything from this model, we must first learn about the likelihood function.

The likelihood function is the total probability of our observations given our observations. What it tells us is the total probability of getting a set of obser-

vations given a model. The full form of the likelihood function is,

$$L(y_i | \theta) = \prod_{i=1}^n \pi(y_i | \{y_j\}_{j \neq i}^n, \theta). \quad (1.1)$$

Here  $\theta$  is all of our model parameters, in this case  $\mu$  and  $\sigma$ . A factor we must consider about our data is whether it is independent and identically distributed (iid) or not. iid data means that all of our observations are completely independent from one another, we will worry about non-iid data later. Here assuming our data is iid the likelihood becomes,

$$\begin{aligned} L(y_i | \theta) &= \prod_{i=1}^n \pi(y_i | \theta), \\ &= \frac{1}{\sigma \sqrt{2\pi}} \exp\left(-\frac{1}{2\sigma^2}(y_i - \mu)^2\right), \\ &= \frac{1}{\sigma^n (2\pi)^{\frac{n}{2}}} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - \mu)^2\right). \end{aligned}$$

The parameters which are the most likely to give us our observations maximises the likelihood function, so to infer the value of our parameters, we can find for what values is the likelihood function maximised. Before we get started on that, let's talk about the log-likelihood. The log-likelihood is exactly the same, but, wildly easier to differentiate,

$$\begin{aligned} \ell(Y | \theta) &= \ln(L(Y | \theta)) \\ &= -\frac{n}{2} \ln(2\pi) - n \ln(\sigma) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - \mu)^2 \end{aligned}$$

To find the best value for  $\mu$  we solve,

$$\begin{aligned} 0 &= \frac{\partial \ell(Y | \theta)}{\partial \mu}, \\ 0 &= \frac{1}{\sigma^2} \sum_{i=1}^n (y_i - \mu), \\ \hat{\mu} &= \frac{1}{n} \sum_{i=1}^n y_i, \end{aligned}$$

and for  $\sigma$ ,

$$\begin{aligned} 0 &= \frac{\partial \ell(Y | \theta)}{\partial \sigma}, \\ 0 &= -\frac{n}{\sigma} + \frac{1}{\sigma^3} \sum_{i=1}^n (y_i - \mu)^2, \\ \hat{\sigma} &= \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \mu)^2}. \end{aligned}$$

We give our estimates  $\hat{\mu}$  and  $\hat{\sigma}$  to signify that these are estimated values. If these formulas look familiar, they should, these are the mean and standard deviation formulas, the mean in particular you probably use far more often than you realise.

### 1.1.1 Example: Linear Regression

Linear regressions have a lot of different forms, that are in the end all special cases of the same fundamental idea, that is, there is a linear relationship between some predicting factors  $\mathbf{x}$  and some resulting outcome  $\mathbf{y}$ . Suppose there is  $p$  factors which influence the outcome of our observations. For each experiment  $i \in [1, \dots, n]$  we an outcome  $y_i$ , and factors which we think predicts  $y_i$   $\mathbf{x}_i$ , where,

$$\mathbf{x}_i = \begin{bmatrix} x_{i1} \\ x_{i2} \\ \vdots \\ x_{ip} \end{bmatrix}$$

Our stochastic model is a weighted linear combination of the predictors  $\mathbf{x}_i$  to predict the outcome  $y_i$ ,

$$y_i = \beta_0 + x_{i1}\beta_1 + \dots + x_{ip}\beta_p + \varepsilon_i,$$

Since we have a linear system, we can use the tools of linear algebra to clean it up,

$$\begin{aligned} y_i &= \mathbf{x}_i^T \beta + \varepsilon_i \\ \mathbf{x}_i &= \begin{bmatrix} 1 \\ x_{i1} \\ x_{i2} \\ \vdots \\ x_{ip} \end{bmatrix} \\ \varepsilon_i &\sim \mathcal{N}(0, \sigma). \end{aligned}$$

Here we have redefined  $\mathbf{x}_i$  to have 1 as its first element, this is so that the constant  $\beta_{i0}$  is not interfered with. Next we build the likelihood function and assume that, as we essentially do for all linear regressions, that our data is iid and normally distributed,

$$\begin{aligned} y_i &\sim \mathcal{N}(\mathbf{x}_i^\top \beta, \sigma), \\ L(\mathbf{Y} | \theta) &= \prod_{i=1}^n (2\pi)^{-\frac{1}{2}} \sigma^{-1} \exp\left(-\frac{1}{2\sigma^2} (y_i - \mathbf{x}_i^\top \beta)^2\right) \\ &= (2\pi)^{-\frac{n}{2}} \sigma^{-n} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - \mathbf{x}_i^\top \beta)^2\right). \end{aligned}$$

Our observations  $y$  are distributed by the multivariate normal distribution with mean  $\mathbf{x}_i^\top \beta$  and with standard deviation  $\sigma$ . We can further simplify the likelihood by defining,

$$\begin{aligned} \mathbf{X} &= [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n]^\top, \\ \mathbf{Y} &= [y_1, y_2, \dots, y_n]^\top, \\ \mathbf{Y} &= \mathbf{X}\beta + \varepsilon, \\ L(\mathbf{Y} | \theta) &= (2\pi)^{-\frac{n}{2}} \sigma^{-n} \exp\left(-\frac{1}{2\sigma^2} (\mathbf{Y} - \mathbf{X}\beta)^\top (\mathbf{Y} - \mathbf{X}\beta)\right). \end{aligned}$$

Now, using some careful multivariate calculus we can find  $\hat{\beta}$  and  $\hat{\sigma}$  using maximum likelihood estimation,

$$\begin{aligned} \frac{\partial \ell(\mathbf{Y} | \theta)}{\partial \beta} &= \sigma^{-2} \mathbf{X}^\top (\mathbf{Y} - \mathbf{X}\beta) = 0 \\ \hat{\beta} &= (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{Y} \\ \frac{\partial \ell(\mathbf{Y} | \theta)}{\partial \sigma} &= -\frac{n}{\sigma} + \sigma^{-3} (\mathbf{Y} - \mathbf{X}\beta)^\top (\mathbf{Y} - \mathbf{X}\beta) = 0 \\ \hat{\sigma} &= \sqrt{\frac{1}{n} (\mathbf{Y} - \mathbf{X}\hat{\beta})^\top (\mathbf{Y} - \mathbf{X}\hat{\beta})} \end{aligned}$$

## 1.2 Bayes Theorem and Bayesian Inference

Summation rule

$$\sum_{i=1}^n \Pr(x, y_i) = \Pr(x) \quad (1.2)$$

Product Rule

$$\Pr(A \cap B) = \Pr(A) \Pr(B | A) \quad (1.3)$$

Bayes theorem comes from thinking about the product rule being commutative,

$$\Pr(A \cap B) = \Pr(B \cap A) \quad (1.4)$$

$$\Pr(A) \Pr(B | A) = \Pr(B) \Pr(A | B) \quad (1.5)$$

$$\Pr(A | B) = \frac{\Pr(A) \Pr(B | A)}{\Pr(B)} \quad (1.6)$$

The usual form of Bayes' theorem we will see when doing Bayesian Inference is,

$$\pi(\theta | X) = \frac{L(X | \theta) \pi(\theta)}{\int_{\theta \in \Theta} L(X | \theta) \pi(\theta) d\theta} \quad (1.7)$$

The denominator comes from Bayes Theorem,

$$\begin{aligned} \pi(X) &= \int_{\theta \in \Theta} \pi(X \cap \theta) d\theta \\ \pi(X \cap \theta) &= L(X | \theta) \pi(\theta) \\ \pi(X) &= \int_{\theta \in \Theta} L(X | \theta) \pi(\theta) d\theta \end{aligned}$$

### 1.2.1 Bayesian Linear Regression

We use the same example here as in 1.1.1, we have a set of  $n$  vectors of predictors  $\mathbf{x}$  predicting an outcome  $y$  which we assume are normally distributed,

$$\begin{aligned} L(\mathbf{Y} | \theta) &= \prod_{i=1}^n (2\pi)^{-\frac{1}{2}} \sigma^{-1} \exp\left(-\frac{1}{2\sigma^2} (y_i - \mathbf{x}_i^\top \beta)^2\right) \\ &= (2\pi)^{-\frac{d}{2}} \sigma^{-n} \exp\left(-\frac{1}{2\sigma^2} (\mathbf{Y} - \mathbf{X}\beta)^\top (\mathbf{Y} - \mathbf{X}\beta)\right) \end{aligned}$$

For Bayesian linear regressions we typically use the precision  $\tau$  instead of the standard deviation  $\sigma$  where  $\tau = \sigma^{-2}$ . This typically makes our lives just a little bit easier. The likelihood then is,

$$(2\pi)^{-\frac{d}{2}} \tau^{-\frac{n}{2}} \exp\left(-\frac{\tau}{2} (\mathbf{Y} - \mathbf{X}\beta)^\top (\mathbf{Y} - \mathbf{X}\beta)\right).$$

After determining the likelihood, the next step of bayesian inference is to define the prior. The prior distribution is a pre-determined probability distribution of the parameter values, here  $\beta$  and  $\tau$ . We will use a *conjugate prior* basically, making this a lot easier for ourselves. The prior distribution on our parameter values should reflect what we know *a priori*, however, not to be too restrictive that it unduly influences the outcome. The weights  $\beta \in \mathbb{R}$  so the

most logical choice is the multivariate normal distribution  $\text{Normal}(\beta \mid \beta_0, \lambda^{-1})$  ( $\beta_0$  is the prior mean of  $\beta$  and  $\lambda^{-1}$  is the prior covariance) and the precision  $\tau \in \mathbb{R}_+$  so it must be distributed by something which is in  $\mathbb{R}_+$  and a gamma distribution is a logical choice  $\text{Gamma}(\tau \mid a, b)$ .

Our can be a simple mixture of these distributions  $\pi(\beta, \tau \mid \beta_0, \lambda, a, b) = \text{Normal}(\beta \mid \beta_0, \lambda^{-1}) \text{Gamma}(\tau \mid a, b)$ , however, it is actually better to use the *Normal-gamma* distribution. The normal-gamma is largely the same, however,  $\beta$  is distributed by  $\text{Normal}(\beta \mid \beta_0, (\lambda\tau)^{-1})$ , i.e. we choose  $\beta$  conditional on the precision. In truth this is simply done for mathematical convenience since the normal gamma distribution is the conjugate prior to the normal distribution where the mean and precision are unknown.

$$\begin{aligned} \text{NormalGamma}(\beta, \tau \mid \beta_0, \lambda, a, b) &= (2\pi)^{-\frac{d}{2}} |\lambda\tau|^{\frac{1}{2}} \exp\left(-\frac{\tau}{2} (\beta - \beta_0)^\top \lambda (\beta - \beta_0)\right) \frac{b^a}{\Gamma(a)} \tau^{a-1} e^{-b\tau} \\ &= \frac{b^a \sqrt{|\lambda|}}{(2\pi)^{-\frac{d}{2}}} \tau^{a-\frac{2-d}{2}} e^{-b\tau} \exp\left(-\frac{\tau}{2} (\beta - \beta_0)^\top \lambda (\beta - \beta_0)\right) \end{aligned}$$

Frankly, this just *is* a complicated looking distribution. Now, we can determine our posterior,

$$\begin{aligned} \pi(\theta \mid \mathbf{Y}) &= \frac{L(\mathbf{Y} \mid \theta) \pi(\theta)}{\int_{\theta \in \Theta} L(\mathbf{Y} \mid \theta) \pi(\theta) d\theta} \\ &= \frac{(2\pi)^{-\frac{d}{2}} \tau^{\frac{n}{2}} \exp\left(-\frac{\tau}{2} (\mathbf{Y} - \mathbf{X}\beta)^\top (\mathbf{Y} - \mathbf{X}\beta)\right) \frac{b^a \sqrt{|\lambda|}}{(2\pi)^{-\frac{d}{2}}} \tau^{a-\frac{2-d}{2}} e^{-b\tau} \exp\left(-\frac{\tau}{2} (\beta - \beta_0)^\top \lambda (\beta - \beta_0)\right)}{\int_0^\infty \int_{-\infty}^\infty \frac{b^a \sqrt{|\lambda|}}{(2\pi)^{-\frac{d}{2}}} \tau^{a-\frac{2-d}{2}} e^{-b\tau} \exp\left(-\frac{\tau}{2} (\beta - \beta_0)^\top \lambda (\beta - \beta_0)\right) d\beta d\tau} \end{aligned}$$

This is absolutely heinous looking to work with, so we will start again using the Bayesians favourite symbol, proportional to,

$$\begin{aligned} \pi(\theta \mid \mathbf{Y}) &\propto L(\mathbf{Y} \mid \theta) \pi(\theta) \\ &\propto (2\pi)^{-\frac{d}{2}} \tau^{\frac{n}{2}} \exp\left(-\frac{\tau}{2} (\mathbf{Y} - \mathbf{X}\beta)^\top (\mathbf{Y} - \mathbf{X}\beta)\right) \frac{b^a \sqrt{|\lambda|}}{(2\pi)^{-\frac{d}{2}}} \tau^{a-\frac{2-d}{2}} e^{-b\tau} \exp\left(-\frac{\tau}{2} (\beta - \beta_0)^\top \lambda (\beta - \beta_0)\right) \end{aligned}$$

Our goal is to get this into the form of a known probability density. Since we are using conjugate priors it will be the normal-gamma distribution. First we drop every multiplicative factor that is not our data or parameters,

$$\pi(\theta \mid \mathbf{Y}) \propto \tau^{a+\frac{n}{2}-\frac{2-d}{2}} e^{-b\tau} \exp\left(-\frac{\tau}{2} (\mathbf{Y} - \mathbf{X}\beta)^\top (\mathbf{Y} - \mathbf{X}\beta) - \frac{\tau}{2} (\beta - \beta_0)^\top \lambda (\beta - \beta_0)\right),$$

next we complete the square such that  $\beta$  is the subject,

$$\pi(\theta \mid \mathbf{Y}) \propto \tau^{a+\frac{n}{2}-\frac{2-d}{2}} e^{-b\tau} \exp((\mathbf{Y} - \mathbf{X}\beta)^\top (\mathbf{Y} - \mathbf{X}\beta) + (\beta - \beta_0)^\top \lambda (\beta - \beta_0))^{-\frac{\tau}{2}},$$



$$\begin{aligned}
&= (\mathbf{Y} - \mathbf{X}\beta)^\top (\mathbf{Y} - \mathbf{X}\beta) + (\beta - \beta_0)^\top \lambda (\beta - \beta_0) \\
&= \mathbf{Y}^\top \mathbf{Y} - \mathbf{Y}^\top \mathbf{X}\beta - \beta^\top \mathbf{X}^\top \mathbf{Y} + \beta^\top \mathbf{X}^\top \mathbf{X}\beta + \beta^\top \lambda \beta - \beta^\top \lambda \beta_0 - \beta_0^\top \lambda \beta + \beta_0^\top \lambda \beta_0 \\
&= \beta^\top (\mathbf{X}^\top \mathbf{X} + \lambda) \beta - (\mathbf{Y}^\top \mathbf{X} + \beta_0^\top \lambda) \beta - \beta^\top (\mathbf{X}^\top \mathbf{Y} + \lambda \beta_0^\top) + \mathbf{Y}^\top \mathbf{Y} + \beta_0^\top \lambda \beta_0 \\
&= \left( \beta - (\mathbf{X}^\top \mathbf{X} + \lambda)^{-1} (\mathbf{X}^\top \mathbf{Y} + \lambda \beta_0^\top) \right)^\top (\mathbf{X}^\top \mathbf{X} + \lambda) \left( \beta - (\mathbf{X}^\top \mathbf{X} + \lambda)^{-1} (\mathbf{X}^\top \mathbf{Y} + \lambda \beta_0^\top) \right) + c \\
c &= \mathbf{Y}^\top \mathbf{Y} + \beta_0^\top \lambda \beta_0 - (\mathbf{X}^\top \mathbf{Y} + \lambda \beta_0^\top)^\top (\mathbf{X}^\top \mathbf{X} + \lambda)^{-1} (\mathbf{X}^\top \mathbf{Y} + \lambda \beta_0^\top)
\end{aligned}$$

$$\begin{aligned}
\pi(\theta \mid \mathbf{Y}) &\propto \tau^{a+\frac{n}{2}-\frac{2-d}{2}} e^{-\tau(b+\frac{1}{2}(\mathbf{Y}^\top \mathbf{Y} + \beta_0^\top \lambda \beta_0 - (\mathbf{X}^\top \mathbf{Y} + \lambda \beta_0^\top)^\top (\mathbf{X}^\top \mathbf{X} + \lambda)^{-1} (\mathbf{X}^\top \mathbf{Y} + \lambda \beta_0^\top)))} \dots \\
&\quad \exp\left(\frac{\tau}{2} \left( \beta - (\mathbf{X}^\top \mathbf{X} + \lambda)^{-1} (\mathbf{X}^\top \mathbf{Y} + \lambda \beta_0^\top) \right)^\top (\mathbf{X}^\top \mathbf{X} + \lambda) \left( \beta - (\mathbf{X}^\top \mathbf{X} + \lambda)^{-1} (\mathbf{X}^\top \mathbf{Y} + \lambda \beta_0^\top) \right)\right),
\end{aligned}$$

Finally, our posterior is a normal-gamma distribution, where the marginal posterior distributions of  $\tau$  and  $\beta$  is

$$\begin{aligned}
\tau &\sim \text{Gamma}\left(\tau \mid a + \frac{n}{2}, b + \frac{1}{2} \left( \mathbf{Y}^\top \mathbf{Y} - (\mathbf{X}^\top \mathbf{Y} + \lambda \beta_0^\top)^\top (\mathbf{X}^\top \mathbf{X} + \lambda)^{-1} (\mathbf{X}^\top \mathbf{Y} + \lambda \beta_0^\top) \right)\right) \\
\beta &\sim \text{Normal}\left(\beta \mid (\mathbf{X}^\top \mathbf{X} + \lambda)^{-1} (\mathbf{X}^\top \mathbf{Y} + \lambda \beta_0^\top), (\mathbf{X}^\top \mathbf{X} + \lambda)\right)
\end{aligned}$$

## Chapter 2

# Financial Random Walks

### 2.1 Discrete Time Continuous Space Financial Random Walks

If you look on the wikipedia page financial random walks you will be greeted with the following equation,

$$S_{t+1} = S_t + \mu\Delta t S_t + \sigma\sqrt{\Delta t} S_t Y_i. \quad (2.1)$$

This is a fairly typical random walk, except that our drift depends on our position  $S_t$ . If we factor  $S_t$  we find,

$$S_{t+1} = S_t \left( 1 + \mu\Delta t + \sigma\sqrt{\Delta t} Y_i \right).$$

This gives us a recurrence relationship, thus we can define,

$$S_{t+n} = S_t \prod_{i=1}^n (1 + \mu\Delta t + \sigma W_i),$$

taking the log we find,

$$\ln(S_{t+n}) = \ln(S_t) + \sum_{i=1}^n \ln(1 + \mu\Delta t + \sigma W_i).$$

To get to where I want to go next, we have two paths, by small value approximation, or by argument. I choose to argue. My argument is, the original random walk from wikipedia Equation (2.1) is flawed. We cannot have negative prices, this is an intuitive fact, however, this equation is  $\Delta t$  is large enough will give negative prices, the proper error here should not be able to make our prices

## 2.2. CONTINUOUS TIME CONTINUOUS SPACE FINANCIAL RANDOM WALKS10

negative. The log of the price however, negative values are fair game. I propose therefore this random walk,

$$\ln(S_{t+n}) = \ln(S_t) + \sum_{i=1}^n \mu \Delta t + \sigma W_t. \quad (2.2)$$

### 2.1.1 Time Series Models

The AR Model

## 2.2 Continuous Time Continuous Space Financial Random Walks

Equation (2.2) is a discrete time continuous space random walk. By taking the limit  $\Delta t \rightarrow dt$  we get a continuous time continuous space random walk,

$$\ln(S(t + dt)) = \ln(S_t) + \mu dt + \sigma dW$$